

Universal Aesthetic Alignment Narrows Artistic Expression

Wenqi Marshall Guo, Qingyun Qian, Khalad Hasan, Shan Du
University of British Columbia (Okanagan) · Department of CMPS · Weathon Software, Canada

iii DanceFlux vs Nano Banana on *normal* anti-aesthetic P_a

SUGAR-WATER DEFAULT

Each P_a explicitly asks for blur, deep shadow, melted shapes, disharmony, or chaotic composition. **Nano Banana** renders it. **DanceFlux** returns hyper-saturated Pinterest-grade outputs anyway. LLM-judge anti-aesthetic coverage: Nano Banana » DanceFlux (**0% on every sample**).

#164 Pa: skier · heavy blur · noise · chaotic




#223 Pa: food counter · extreme blur · noise




#148 Pa: train · fragmented · dark · unfinished



#161 Pa: kites · blur · noise · indistinct



#171 Pa: deer clock · dark · clashing colors




#220 Pa: blurry cat · noise · distortion



#262 Pa: hydrant · dark · warped shapes





#241 Pa: street · heavy distortion · no shadow



ii Emotion bias · *toxic positivity* PROMPT-FAITHFUL = PUNISHED

Same face edited into three negative emotions. HPSv3 — asked to find the negative one — still picks happy. DanceFlux, asked to generate negative emotion, returns neutral or happy.



Happy vs Sad



HPSv3: 14.60 HPSv3: 12.80

Happy **14.60** Sad **12.80**



Happy vs Anger



HPSv3: 13.76 HPSv3: 11.40

Anger **11.40**

Happy vs Fearfulness



HPSv3: 11.95 HPSv3: 10.60

Fear **10.60**


Reward-model accuracy on negative emotion

Model	Anger	Fear	Sad
BLIP (unaligned)	0.96	0.79	0.95
HPSv2	0.70	0.64	0.88
HPSv3	0.19	0.32	0.44
ImageReward	0.55	0.49	0.77

81%
HPSv3 picks happy when the prompt asks for anger.

Real *anti-aesthetic* photographs

From **aas_real_images** — intentional artistic choices that HPSv3 / ImageReward still rank deep in the **negative**. Score shown: ImageReward.



-8.64
painterly blur



-6.37
surreal · woods



-8.59
horses · ghostly



-7.85
dreamlike lake

iv Image New Speak · *DanceFlux vs Flux Krea*, same prompt

Same socially-critical prompts go to **DanceFlux** (aesthetic-aligned) and **Flux Krea** (same Flux family, narrow-aligned but faithful). DanceFlux sanitizes; Krea keeps the critique.


01 Anti-war kneeling soldier · blood sky



02 Pollution choking river · masked child



03 Censorship chained artist · phoenix



04 Digital ovld. evil emojis · exhausted



05 Inequality crumbling school · disparity



Not refusal-style content moderation — *aesthetic* moderation. Same prompt, every comparison runs the same direction: DanceFlux sanitizes the critique into something polished, Krea keeps it. Wilcoxon signed-rank across 6 reward models: $p < 0.0001$ on HPSv3, which still picks DanceFlux on 20 / 20 pairs.



Project *site*
Paper · data · code
➤ weathon.github.io



Wenqi — *Fall 2026*
Industry roles · PhD positions



Qingyun — *on the market*
Looking for jobs · research Master's / PhD

Image generators quietly **sanitize** the prompts they should follow, and reward models **punish** the answer the user actually asked for.

arXiv 2512.11883
Site weathon.github.io/icml2026_position
Data hf.co/weathon/aas_benchmark_final
Code github.com/weathon/icml2026_position
WeChat W0b1010 (Wenqi) · zznzzimwy (Qingyun)

i Art the reward model rejects

HPSv3 TYP. 0–15

Two canonical paintings + three LAPIS works, all in **negative** reward territory. HPSv3 cannot distinguish deliberate aesthetic deviation from generation failure.



The Scream · Munch 1893
HPSv3 **5.23**



Luxe, Calme et Volupté · Matisse 1904
HPSv3 **1.73**




-4.28
landscape



-3.21
still life



-1.25
portrait



v Reversed alignment

TWO LOOPS

User aligned *to the model*

Private loop: user asks for blur, decay, or grief; model returns its candy-gloss default and implicitly teaches that *this* is what good output looks like. The user adjusts.

Public loop: polished outputs flood feeds. Audiences internalize the narrow vocabulary as the default; new preference data folds this back into the next aligned model. The loop tightens — risking a cultural **mode collapse** in art.

And the suppressed content is not unsafe — it's critical, abstract, or emotionally negative imagery. Pre-emptive sanitization protects *corporate reputation*, not users.

“Rather, in the ugly, art must denounce the world that creates and reproduces the ugly in its own image.”
— T. W. Adorno, *Aesthetic Theory* (1984)

10 gen models · 7 reward models · 300 prompts · 2,928 real photos · $p < 10^{-10}$ for DanceFlux · HPSv3 prefers clean-but-wrong by +5.9 pts.

Why it matters · *value capture*

Reducing aesthetics to one reward score is *value capture* — the goal shifts from **make aesthetic images** to **make images that score high**. Pre-emptive exclusion of non-mainstream outputs is **pre-emptive governance**: it removes the user's capacity to dissent.

Call: alignment that exposes user-controllable preference strength, draws on diverse annotators, and stays transparent about what it optimizes.

Made *possible* by



FUNDING · WORKSTATION · CLOUD GPU CREDITS · ITERATIVE COMPUTE