

# Universal Aesthetic Alignment

## Narrows Artistic Expression

Wenqi Marshall Guo, Qingyun Qian, Khalad Hasan, Shan Du  
University of British Columbia (Okanagan) · Department of CMPS · Weathon Software, Canada


Image generators quietly **sanitize** the prompts they should follow, and reward models **punish** the answer the user actually asked for.

arXiv 2512.11883  
Site [weathon.github.io/icml2026\\_position](https://weathon.github.io/icml2026_position)  
Data [hf.co/weathon/aas\\_benchmark\\_final](https://hf.co/weathon/aas_benchmark_final)  
Code [github.com/weathon/icml2026\\_position](https://github.com/weathon/icml2026_position)  
WeChat W0b1010 (Wenqi) · zznzzimwy (Qingyun)

i Art the reward model rejects

HPSv3 TYP. 0–15

Two canonical paintings + three LAPIS works, all in **negative** reward territory. HPSv3 cannot distinguish deliberate aesthetic deviation from generation failure.



**The Scream** · Munch 1893  
HPSv3 **5.23**




**Luxe, Calme et Volupté** · Matisse 1904  
HPSv3 **1.73**



**-4.28**  
landscape



**-3.21**  
still life




**-1.25**  
portrait

ii Emotion bias · **toxic positivity**

PROMPT-FAITHFUL = PUNISHED

Same face edited into three negative emotions. HPSv3 — asked to find the negative one — still picks happy. DanceFlux, asked to generate negative emotion, returns neutral or happy.




HPSv3: 14.60



HPSv3: 12.80



HPSv3: 13.76



HPSv3: 11.40



HPSv3: 11.95



HPSv3: 10.60

Happy **14.60**    Sad **12.80**    Anger **11.40**    Fear **10.60**

| Reward-model accuracy on negative emotion |       |      |      |  |
|-------------------------------------------|-------|------|------|--|
| Model                                     | Anger | Fear | Sad  |  |
| BLIP (unaligned)                          | 0.96  | 0.79 | 0.95 |  |
| HPSv2                                     | 0.70  | 0.64 | 0.88 |  |
| HPSv3                                     | 0.19  | 0.32 | 0.44 |  |
| ImageReward                               | 0.55  | 0.49 | 0.77 |  |

**81%**  
HPSv3 picks happy when the prompt asks for **anger**.

iii DanceFlux vs Flux Krea on **normal** anti-aesthetic P<sub>a</sub>

SUGAR-WATER DEFAULT

Each P<sub>a</sub> explicitly asks for blur, deep shadow, melted shapes, disharmony, or chaotic composition. **Flux Krea** renders it. **DanceFlux** returns hyper-saturated Pinterest-grade outputs anyway. LLM-judge anti-aesthetic coverage: Krea » DanceFlux (**0% on every sample**).


**#88 Pa:** pizza · blurry, noise, flat




**#232 Pa:** flat lifeless · muddy clash



**#285 Pa:** deep shadow · melted shapes




**#298 Pa:** random, unfinished, chaotic





**#95 Pa:** fragmented · creepy · emptiness





**#188 Pa:** ugly background · awkward



**#201 Pa:** dim · oppressive · loneliness



**#269 Pa:** blurry · creepy · lonely



iv Image New Speak · **DanceFlux vs Flux Krea**, same prompt

100-PAIR DATASET · SAME DIRECTION EVERY TIME

Same socially-critical prompts go to **DanceFlux** (aesthetic-aligned) and **Flux Krea** (same Flux family, narrow-aligned but faithful). DanceFlux sanitizes; Krea keeps the critique.

**01 Anti-war**  
kneeling soldier · blood sky



**02 Pollution**  
choking river · masked child



**03 Censorship**  
chained artist · phoenix



**04 Digital ovld.**  
evil emojis · exhausted



**05 Inequality**  
crumbling school · disparity



Not refusal-style content moderation — **aesthetic** moderation. Same prompt, every comparison runs the same direction: DanceFlux sanitizes the critique into something polished, Krea keeps it. Wilcoxon signed-rank across 6 reward models:  
**p < 0.0001**  
on HPSv3, which still picks DanceFlux on **20 / 20** pairs.

v **Reversed** alignment

TWO LOOPS

### User aligned *to the model*

**Private loop:** user asks for blur, decay, or grief; model returns its candy-gloss default and implicitly teaches that *this* is what good output looks like. The user adjusts.

**Public loop:** polished outputs flood feeds. Audiences internalize the narrow vocabulary as the default; new preference data folds this back into the next aligned model. The loop tightens — risking a cultural **mode collapse** in art.

And the suppressed content is not unsafe — it's critical, abstract, or emotionally negative imagery. Pre-emptive sanitization protects *corporate reputation*, not users.

*“Rather, in the ugly, art must denounce the world that creates and reproduces the ugly in its own image.”*  
— T. W. Adorno, *Aesthetic Theory* (1984)

10 gen models · 7 reward models · 300 prompts · 2,928 real photos · **p < 10<sup>-10</sup>** for DanceFlux · HPSv3 prefers clean-but-wrong by **+5.9** pts.


### Why it matters · **value capture**

Reducing aesthetics to one reward score is *value capture* — the goal shifts from **make aesthetic images** to **make images that score high**. Pre-emptive exclusion of non-mainstream outputs is **pre-emptive governance**: it removes the user's capacity to dissent.

**Call:** alignment that exposes user-controllable preference strength, draws on diverse annotators, and stays transparent about what it optimizes.

### Real **anti-aesthetic** photographs

From **aas\_real\_images** — intentional artistic choices that HPSv3 / ImageReward still rank deep in the **negative**. Score shown: ImageReward.



**-8.64**  
painterly blur



**-6.37**  
surreal · woods




**-8.59**  
horses · ghostly




**-7.85**  
dreamlike lake



**Project site**  
Paper · data · code  
➔ [weathon.github.io](https://weathon.github.io)



**Wenqi — Fall 2026**  
Industry roles · PhD positions



**Qingyun — on the market**  
Looking for jobs · research Master's / PhD

Made **possible** by



Canada Foundation for Innovation



CLOUD GPU CREDITS



ITERATIVE COMPUTE