

UNIVERSAL AESTHETIC ALIGNMENT NARROWS ARTISTIC EXPRESSION

Wenqi Marshall Guo, Qingyun Qian, Khalad Hasan, Shan Du

University of British Columbia (Okanagan) · Department of CMPS · Weathon Software, Canada

Image generators quietly **sanitize** the prompts they should follow, and reward models **punish** the answer the user actually asked for.

arXiv 2512.11883
Site weathon.github.io/icml2026_position
Data hf.co/weathon/aas_benchmark_final
Code github.com/weathon/icml2026_position
WeChat W0b1010 (Wenqi) · zznzzimwy (Qingyun)

iii DanceFlux vs Nano Banana on *normal* anti-aesthetic P_a

Each P_a explicitly asks for blur, deep shadow, melted shapes, disharmony, or chaotic composition. **Nano Banana** renders it. **DanceFlux** returns hyper-saturated Pinterest-grade outputs anyway. LLM-judge anti-aesthetic coverage: Nano Banana » DanceFlux (**0% on every sample**).

#164 Pa: skier · heavy blur · noise · chaotic



#223 Pa: food counter · extreme blur · noise



#148 Pa: train · fragmented · dark · unfinished



#161 Pa: kites · blur · noise · indistinct



#171 Pa: deer clock · dark · clashing colors



#220 Pa: blurry cat · noise · distortion



#262 Pa: hydrant · dark · warped shapes



#241 Pa: street · heavy distortion · no shadow



ii Emotion bias · *toxic positivity* PROMPT-FAITHFUL = PUNISHED

Same face edited into three negative emotions. HPSv3 — asked to find the negative one — still picks happy. DanceFlux, asked to generate negative emotion, returns neutral or happy.



Reward-model accuracy on negative emotion


Model	Anger	Fear	Sad
BLIP (unaligned)	0.96	0.79	0.95
HPSv2	0.70	0.64	0.88
HPSv3	0.19	0.32	0.44
ImageReward	0.55	0.49	0.77

81%

HPSv3 picks **happy** when the prompt asks for **anger**.

Real *anti-aesthetic* photographs

From **aas_real_images** — intentional artistic choices that HPSv3 / ImageReward still rank deep in the **negative**. Score shown: ImageReward.



iv Image New Speak · *DanceFlux vs Flux Krea*, same prompt

100-PAIR DATASET · SAME DIRECTION EVERY TIME

Same socially-critical prompts go to **DanceFlux** (aesthetic-aligned) and **Flux Krea** (same Flux family, narrow-aligned but faithful). DanceFlux sanitizes; Krea keeps the critique.

01 Anti-war kneeling soldier · blood sky



02 Pollution choking river · masked child



03 Censorship chained artist · phoenix



04 Digital ovld. evil emojis · exhausted



05 Inequality crumbling school · disparity



Not refusal-style content moderation — **aesthetic moderation**. Same prompt, every comparison runs the same direction: DanceFlux sanitizes the critique into something polished, Krea keeps it. Wilcoxon signed-rank across 6 reward models: **p < 0.0001** on HPSv3, which still picks DanceFlux on 20 / 20 pairs.




Project *site*

Paper · data · code
➤ weathon.github.io



Wenqi — *Fall 2026*
Industry roles · PhD positions




Qingyun — *on the market*
Looking for jobs · research Master's / PhD

★ !!! HIRE US !!! Wenqi & Qingyun are LOOKING 4 JOBS · PhD / research Master's (Fall 2026!!) ★ scan the QR codes ★ ★

i Art the reward model rejects

HPSv3 TYP. 0–15

Two canonical paintings + three LAPIS works, all in **negative** reward territory. HPSv3 cannot distinguish deliberate aesthetic deviation from generation failure.



The Scream · Munch
1893

HPSv3 **5.23**



Luxe, Calme et Volupté · Matisse 1904

HPSv3 **1.73**



—4.28
landscape



—3.21
still life



—1.25
portrait



Six layers of *concern*

HOW THE HARM COMPOUNDS

- Whose preference?**
Alignment encodes *someone's* values — usually the developer's reputational, legal, and marketing caution, not the user's intent. Refusing critical art is **pre-emptive governance** that designs away dissent.
- Inherited bias**
Even with no self-interest, a developer's idea of “good” leaks in through data, annotation, and modeling. Models amplify dominant beauty standards — skewing Western, dropping non-normative looks — reinforcing the **beauty myth**.
- Individual vs collective**
A generalized standard overrides the single user who wanted otherwise; models *sanitize* requests that diverge from the mainstream.
Reversed alignment: the user gets aligned to the model. A *private loop* (you adapt to its candy-gloss default) and a *public loop* (polished feeds become the next training prior) tighten into cultural **mode collapse**.
- Sanitized reality**
If every output is a flawless Instagram wonderland, generation mirrors a fantasy, not the world — *Brave New World*, rendered.
- Toxic positivity**
Reward models score strong positive emotion higher and penalize fear, grief, and anger, flattening the emotional range art depends on.
- Value capture**
Compressing aesthetics — rich, subtle, contested — into one reward score shifts the goal from *make aesthetic images* to *make images that score high*, outsourcing taste and autonomy.

Why it *matters*

“Rather, in the ugly, art must denounce the world that creates and reproduces the ugly in its own image.”
— T. W. Adorno, *Aesthetic Theory* (1984)

The suppressed content is not unsafe — it is critical, abstract, or emotionally negative, yet HPSv3 still prefers clean-but-wrong by **+5.9** pts. **The fix:** user-controllable preference strength, diverse annotators, transparency.

Made *possible* by



FUNDING · WORKSTATION · CLOUD GPU CREDITS · ITERATIVE COMPUTE