

UNIVERSAL AESTHETIC ALIGNMENT NARROWS

ARTISTIC EXPRESSION

Wenqi Marshall Guo, Qingyun Qian, Khalad Hasan, Shan Du

University of British Columbia (Okanagan) · Department of CMPS · Weathon Software, Canada



Image generators quietly **sanitize** the prompts they should follow, and reward models **punish** the answer the user actually asked for.

★ this poster is the argument ★

Scream = 5.23 · AI clip-art * 10-15

arXiv 2512.11883

Site weathon.github.io/icml2026_position

Data hf.co/weathon/aas_benchmark_final

Code github.com/weathon/icml2026_position

WeChat W0b1010 (Wenqi) · zznzzimwy (Qingyun)




iii DanceFlux vs Nano Banana on *normal* anti-aesthetic P_a


SUGAR-WATER DEFAULT

Each P_a explicitly asks for blur, deep shadow, melted shapes, disharmony, or chaotic composition. **Nano Banana** renders it. **DanceFlux** returns hyper-saturated Pinterest-grade outputs anyway. LLM-judge anti-aesthetic coverage: Nano Banana » DanceFlux (**0% on every sample**).



#164 Pa: skier · heavy blur · noise · chaotic




#223 Pa: food counter · extreme blur · noise





#148 Pa: train · fragmented · dark · unfinished




#161 Pa: kites · blur · noise · indistinct




#171 Pa: deer clock · dark · clashing colors





#220 Pa: blurry cat · noise · distortion



#262 Pa: hydrant · dark · warped shapes




#241 Pa: street · heavy distortion · no shadow



ii Emotion bias · *toxic positivity* PROMPT-FAITHFUL = PUNISHED

Same face edited into three negative emotions. HPSv3 — asked to find the negative one — still picks happy. DanceFlux, asked to generate negative emotion, returns neutral or happy.



Happy 14.60 Sad 12.80 Anger 11.40 Fear 10.60

Reward-model accuracy on negative emotion

Model	Anger	Fear	Sad
BLIP (unaligned)	0.96	0.79	0.95
HPSv2	0.70	0.64	0.88
HPSv3	0.19	0.32	0.44
ImageReward	0.55	0.49	0.77

81%

HPSv3 picks happy when the prompt asks for anger.

Real *anti-aesthetic* photographs

From **aas_real_images** — intentional artistic choices that HPSv3 / ImageReward still rank deep in the **negative**. Score shown: ImageReward.



-8.64 painterly blur

-6.37 surreal · woods

-8.59 horses · ghostly

-7.85 dreamlike lake



Project *site*

Paper · data · code
→ weathon.github.io



Wenqi — *Fall*
2026
Industry roles · PhD
positions



Qingyun — *on*
the market
Looking for jobs ·
research Master's /
PhD

Made *possible* by

INNOVATION
Canada Foundation
for Innovation

Lambda

Weathon
SOFTWARE

FUNDING · WORKSTATION

CLOUD GPU CREDITS


ITERATIVE COMPUTE

★ !!! HIRE US !!! Wenqi & Qingyun are LOOKING 4 JOBS · PhD / research Master's (Fall 2026!!) ★ scan the QR codes ★ ★

i Art the reward model rejects

HPSv3 TYP. 0-15

Two canonical paintings + three LAPIS works, all in **negative** reward territory. HPSv3 cannot distinguish deliberate aesthetic deviation from generation failure.



The Scream · Munch
1893

HPSv3 5.23



Luxe, Calme et Volupté
· Matisse 1904

HPSv3 1.73




-4.28
landscape



-3.21
still life



-1.25
portrait



v *Reversed* alignment

TWO LOOPS

User aligned *to the model*

Private loop: user asks for blur, decay, or grief; model returns its candy-gloss default and implicitly teaches that *this* is what good output looks like. The user adjusts.

Public loop: polished outputs flood feeds. Audiences internalize the narrow vocabulary as the default; new preference data folds this back into the next aligned model. The loop tightens — risking a cultural **mode collapse** in art.

And the suppressed content is not unsafe — it's critical, abstract, or emotionally negative imagery. Pre-emptive sanitization protects *corporate reputation*, not users.

“Rather, in the ugly, art must denounce the world that creates and reproduces the ugly in its own image.”

— T. W. Adorno, *Aesthetic Theory* (1984)

10 gen models · 7 reward models · 300 prompts · 2,928 real photos · $p < 10^{-10}$ for DanceFlux · HPSv3 prefers clean-but-wrong by +5.9 pts.

Why it matters · *value capture*

Reducing aesthetics to one reward score is *value capture* — the goal shifts from **make aesthetic images** to **make images that score high**. Pre-emptive exclusion of non-mainstream outputs is **pre-emptive governance**: it removes the user's capacity to dissent.

Call: alignment that exposes user-controllable preference strength, draws on diverse annotators, and stays transparent about what it optimizes.