

Universal Aesthetic Alignment Narrows Artistic Expression

Wenqi Marshall Guo, Qingyun Qian, Khalad Hasan, Shan Du
University of British Columbia (Okanagan) · Department of CMPS · Weathon Software, Canada

iii DanceFlux vs Nano Banana on *normal* anti-aesthetic P_a

SUGAR-WATER DEFAULT

Each P_a explicitly asks for blur, deep shadow, melted shapes, disharmony, or chaotic composition. **Nano Banana** renders it. **DanceFlux** returns hyper-saturated Pinterest-grade outputs anyway. LLM-judge anti-aesthetic coverage: Nano Banana » DanceFlux (**0% on every sample**).

#164 Pa: skier · heavy blur · noise · chaotic




#223 Pa: food counter · extreme blur · noise




#148 Pa: train · fragmented · dark · unfinished



#161 Pa: kites · blur · noise · indistinct



#171 Pa: deer clock · dark · clashing colors




#220 Pa: blurry cat · noise · distortion



#262 Pa: hydrant · dark · warped shapes





#241 Pa: street · heavy distortion · no shadow




ii Emotion bias · *toxic positivity* PROMPT-FAITHFUL = PUNISHED

Same face edited into three negative emotions. HPSv3 — asked to find the negative one — still picks happy. DanceFlux, asked to generate negative emotion, returns neutral or happy.



Happy vs Sad



Happy vs Anger



Happy vs Fearfulness



Happy 14.60

Sad 12.80

Anger 11.40

Fear 10.60


Reward-model accuracy on negative emotion

Model	Anger	Fear	Sad
BLIP (unaligned)	0.96	0.79	0.95
HPSv2	0.70	0.64	0.88
HPSv3	0.19	0.32	0.44
ImageReward	0.55	0.49	0.77

81%
HPSv3 picks happy when the prompt asks for anger.

Real *anti-aesthetic* photographs

From **aas_real_images** — intentional artistic choices that HPSv3 / ImageReward still rank deep in the **negative**. Score shown: ImageReward.



-8.64
painterly blur



-6.37
surreal · woods



-8.59
horses · ghostly



-7.85
dreamlike lake

iv Image New Speak · *DanceFlux vs Flux Krea*, same prompt

100-PAIR DATASET · SAME DIRECTION EVERY TIME

Same socially-critical prompts go to **DanceFlux** (aesthetic-aligned) and **Flux Krea** (same Flux family, narrow-aligned but faithful). DanceFlux sanitizes; Krea keeps the critique.

01 Anti-war kneeling soldier · blood sky



02 Pollution choking river · masked child



03 Censorship chained artist · phoenix



04 Digital ovld. evil emojis · exhausted



05 Inequality crumbling school · disparity



Not refusal-style content moderation — *aesthetic* moderation. Same prompt, every comparison runs the same direction: DanceFlux sanitizes the critique into something polished, Krea keeps it. Wilcoxon signed-rank across 6 reward models:
 $p < 0.0001$ on HPSv3, which still picks DanceFlux on 20 / 20 pairs.


◆ Image generators quietly **sanitize** the prompts they should follow, and reward models **punish** the answer the user actually asked for.

arXiv 2512.11883
Site weathon.github.io/icml2026_position
Data hf.co/weathon/aas_benchmark_final
Code github.com/weathon/icml2026_position
WeChat W0b1010 (Wenqi) · zznzzimwy (Qingyun)

i Art the reward model rejects

HPSv3 TYP. 0–15


Two canonical paintings + three LAPIS works, all in **negative** reward territory. HPSv3 cannot distinguish deliberate aesthetic deviation from generation failure.




The Scream · Munch 1893
HPSv3 5.23




Luxe, Calme et Volupté · Matisse 1904
HPSv3 1.73



-4.28
landscape



-3.21
still life



-1.25
portrait



v *Reversed* alignment TWO LOOPS

User aligned *to the model*

Private loop: user asks for blur, decay, or grief; model returns its candy-gloss default and implicitly teaches that *this* is what good output looks like. The user adjusts.

Public loop: polished outputs flood feeds. Audiences internalize the narrow vocabulary as the default; new preference data folds this back into the next aligned model. The loop tightens — risking a cultural **mode collapse** in art.

And the suppressed content is not unsafe — it's critical, abstract, or emotionally negative imagery. Pre-emptive sanitization protects *corporate reputation*, not users.

“Rather, in the ugly, art must denounce the world that creates and reproduces the ugly in its own image.”
— T. W. Adorno, *Aesthetic Theory* (1984)

10 gen models · 7 reward models · 300 prompts · 2,928 real photos · $p < 10^{-10}$ for DanceFlux · HPSv3 prefers clean-but-wrong by +5.9 pts.

Why it matters · *value capture*

Reducing aesthetics to one reward score is *value capture* — the goal shifts from **make aesthetic images** to **make images that score high**. Pre-emptive exclusion of non-mainstream outputs is **pre-emptive governance**: it removes the user's capacity to dissent.

Call: alignment that exposes user-controllable preference strength, draws on diverse annotators, and stays transparent about what it optimizes.

Made *possible* by



FUNDING · WORKSTATION · CLOUD GPU CREDITS · ITERATIVE COMPUTE



Project *site*
Paper · data · code
➤ weathon.github.io



Wenqi — *Fall 2026*
Industry roles · PhD positions



Qingyun — *on the market*
Looking for jobs · research Master's / PhD