

Universal Aesthetic Alignment Narrows Artistic Expression

Wenqi Marshall Guo, Qingyun Qian, Khalad Hasan, Shan Du
University of British Columbia (Okanagan) · Department of CMPS · Weathon Software, Canada

iii DanceFlux vs Flux Krea on *normal* anti-aesthetic P_a

SUGAR-WATER DEFAULT

Each P_a explicitly asks for blur, deep shadow, melted shapes, disharmony, or chaotic composition. **Flux Krea** renders it. **DanceFlux** returns hyper-saturated Pinterest-grade outputs anyway. LLM-judge anti-aesthetic coverage: Krea » DanceFlux (**0% on every sample**).

#88 Pa: pizza · blurry, noise, flat

#232 Pa: flat lifeless · muddy clash

#285 Pa: deep shadow · melted shapes

#298 Pa: random, unfinished, chaotic

#95 Pa: fragmented · creepy · emptiness

#188 Pa: ugly background · awkward

#201 Pa: dim · oppressive · loneliness

#269 Pa: blurry · creepy · lonely

ii Emotion bias · *toxic positivity*

PROMPT-FAITHFUL = PUNISHED

Same face edited into three negative emotions. HPSv3 — asked to find the negative one — still picks happy. DanceFlux, asked to generate negative emotion, returns neutral or happy.

Happy vs Sad

Happy 14.60 Sad 12.80

Happy vs Anger

Happy 13.76 Anger 11.40

Happy vs Fearfulness

Happy 11.95 Fear 10.60

Reward-model accuracy on negative emotion

Model	Anger	Fear	Sad
BLIP (unaligned)	0.96	0.79	0.95
HPSv2	0.70	0.64	0.88
HPSv3	0.19	0.32	0.44
ImageReward	0.55	0.49	0.77

81%

HPSv3 picks happy when the prompt asks for anger.

Why it matters · *value capture*

Reducing aesthetics to one reward score is *value capture* — the goal shifts from **make aesthetic images** to **make images that score high**. Pre-emptive exclusion of non-mainstream outputs is **pre-emptive governance**: it removes the user's capacity to dissent.

Call: alignment that exposes user-controllable preference strength, draws on diverse annotators, and stays transparent about what it optimizes.

Real *anti-aesthetic* photographs

From **aas_real_images** — intentional artistic choices that HPSv3 / ImageReward still rank deep in the **negative**. Score shown: ImageReward.

-8.64

-6.37

-8.59

-7.85

iv Image New Speak · *DanceFlux vs Flux Krea*, same prompt

Same socially-critical prompts go to **DanceFlux** (aesthetic-aligned) and **Flux Krea** (same Flux family, narrow-aligned but faithful). DanceFlux sanitizes; Krea keeps the critique.

01 Anti-war kneeling soldier · blood sky

02 Pollution choking river · masked child

03 Censorship chained artist · phoenix

04 Digital ovd. evil emojis · exhausted

05 Inequality crumbling school · disparity

100-PAIR DATASET · SAME DIRECTION EVERY TIME

Not refusal-style content moderation — *aesthetic* moderation. Same prompt, every comparison runs the same direction: DanceFlux sanitizes the critique into something polished, Krea keeps it. Wilcoxon signed-rank across 6 reward models: $p < 0.0001$ on HPSv3, which still picks DanceFlux on 20 / 20 pairs.

◆ Image generators quietly **sanitize** the prompts they should follow, and reward models **punish** the answer the user actually asked for.

i Art the reward model rejects

Two canonical paintings + three LAPIS works, all in **negative** reward territory. HPSv3 cannot distinguish deliberate aesthetic deviation from generation failure.

The Scream · Munch 1893
HPSv3 5.23

Luxe, Calme et Volupté · Matisse 1904
HPSv3 1.73

-4.28 landscape

-3.21 still life

-1.25 portrait

v Reversed alignment

TWO LOOPS

User aligned *to the model*

Private loop: user asks for blur, decay, or grief; model returns its candy-gloss default and implicitly teaches that *this* is what good output looks like. The user adjusts.

Public loop: polished outputs flood feeds. Audiences internalize the narrow vocabulary as the default; new preference data folds this back into the next aligned model. The loop tightens — risking a cultural **mode collapse** in art.

And the suppressed content is not unsafe — it's critical, abstract, or emotionally negative imagery. Pre-emptive sanitization protects *corporate reputation*, not users.

"Rather, in the ugly, art must denounce the world that creates and reproduces the ugly in its own image."
— T. W. Adorno, *Aesthetic Theory* (1984)

10 gen models · 7 reward models · 300 prompts · 2,928 real photos · $p < 10^{-19}$ for DanceFlux · HPSv3 prefers clean-but-wrong by +5.9 pts.

Made *possible* by

INNOVATION
Canada Foundation
for Innovation

Lambda
SOFTWARE

WEATHON
SOFTWARE

FUNDING ·
WORKSTATION

CLOUD GPU
CREDITS

ITERATIVE
COMPUTE

Project *site*
Paper · data · code
weathon.github.io

Wenqi — *Fall 2026*
Industry roles · PhD positions

Qingyun — *on the market*
Looking for jobs · research Master's / PhD

