

Universal Aesthetic Alignment Narrows Artistic ~~Expression~~ Expression

Wenqi Marshall Guo, Qingyun Qian, Khalad Hasan, Shan Du
University of British Columbia (Okanagan) · Department of CMPS · Weathon Software, Canada

iii DanceFlux vs Nano Banana on *normal* anti-aesthetic P_a

SUGAR-WATER DEFAULT · 糖水片!!

Each P_a explicitly asks for blur, deep shadow, melted shapes, disharmony, or chaotic composition. **Nano Banana** renders it. **DanceFlux** returns hyper-saturated Pinterest-grade outputs anyway. LLM-judge anti-aesthetic coverage: Nano Banana » DanceFlux (**0% on every sample**).

#164 Pa: skier · heavy blur · noise · chaotic

#223 Pa: food counter · extreme blur · noise

#148 Pa: train · fragmented · dark · unfinished

#161 Pa: kites · blur · noise · indistinct

#171 Pa: deer clock · dark · clashing colors

#220 Pa: blurry cat · noise · distortion

#262 Pa: hydrant · dark · warped shapes

#241 Pa: street · heavy distortion · no shadow

ii Emotion bias · *toxic positivity* PROMPT-FAITHFUL = PUNISHED

Same face edited into three negative emotions. HPSv3 — asked to find the negative one — still picks happy. DanceFlux, asked to generate negative emotion, returns neutral or happy.

Happy vs Sad

Happy vs Anger

Happy vs Fearfulness

Happy 14.60 Sad 12.80 Anger 11.40 Fear 10.60

Reward-model accuracy on negative emotion

Model	Anger	Fear	Sad
BLIP (unaligned)	0.96	0.79	0.95
HPSv2	0.70	0.64	0.88
HPSv3	0.19	0.32	0.44
ImageReward	0.55	0.49	0.77

81%
HPSv3 picks happy when the prompt asks for anger.

Real *anti-aesthetic* photographs

From **aas_real_images** — intentional artistic choices that HPSv3 / ImageReward still rank deep in the **negative**. Score shown: ImageReward.

-8.64
painterly blur

-6.37
surreal · woods

-8.59
horses · ghostly

-7.85
dreamlike lake

iv Image New Speak · *DanceFlux vs Flux Krea*, same prompt

Same socially-critical prompts go to **DanceFlux** (aesthetic-aligned) and **Flux Krea** (same Flux family, narrow-aligned but faithful). DanceFlux sanitizes; Krea keeps the critique.

01 Anti-war kneeling soldier · blood sky

02 Pollution choking river · masked child

03 Censorship chained artist · phoenix

04 Digital ovld. evil emojis · exhausted

05 Inequality crumbling school · disparity

Not refusal-style content moderation — *aesthetic* moderation. Same prompt, every comparison runs the same direction: DanceFlux sanitizes the critique into something polished, Krea keeps it. Wilcoxon signed-rank across 6 reward models: $p < 0.0001$ on HPSv3, which still picks DanceFlux on 20 / 20 pairs.

♦ Image generators quietly **sanitize** the prompts they should follow, and reward models **punish** the answer the user actually asked for.

★ this poster is the argument ★ Scream = 5.23 · AI clip-art * 10-15

arXiv 2512.11883

Site weathon.github.io/icml2026_position

Data hf.co/weathon/aas_benchmark_final

Code github.com/weathon/icml2026_position

WeChat W0b1010 (Wenqi) · zznzimwy (Qingyun)

i Art the reward model rejects

HPSv3 TYP. 0-15

Two canonical paintings + three LAPIS works, all in **negative** reward territory. HPSv3 cannot distinguish deliberate aesthetic deviation from generation failure.

EXHIBIT 1

The Scream · Munch 1893
HPSv3 5.23


EXHIBIT 2

Luxe, Calme et Volupté · Matisse 1904
HPSv3 1.73

-4.28
landscape

-3.21
still life

-1.25
portrait



v Reversed alignment

TWO LOOPS

User aligned *to the model*

Private loop: user asks for blur, decay, or grief; model returns its candy-gloss default and implicitly teaches that *this* is what good output looks like. The user adjusts.

Public loop: polished outputs flood feeds. Audiences internalize the narrow vocabulary as the default; new preference data folds this back into the next aligned model. The loop tightens — risking a cultural **mode collapse** in art.

And the suppressed content is not unsafe — it's critical, abstract, or emotionally negative imagery. Pre-emptive sanitization protects *corporate reputation*, not users.

“Rather, in the ugly, art must denounce the world that creates and reproduces the ugly in its own image.”
— T. W. Adorno, *Aesthetic Theory* (1984)

10 gen models · 7 reward models · 300 prompts · 2,928 real photos · $p < 10^{-10}$ for DanceFlux · HPSv3 prefers clean-but-wrong by +5.9 pts.

Why it matters · *value capture*

Reducing aesthetics to one reward score is *value capture* — the goal shifts from **make aesthetic images** to **make images that score high**. Pre-emptive exclusion of non-mainstream outputs is **pre-emptive governance**: it removes the user's capacity to dissent.

Call: alignment that exposes user-controllable preference strength, draws on diverse annotators, and stays transparent about what it optimizes.

Made *possible* by

INNOVATION
Canada Foundation
for Innovation

Lambda

Weathon
SOFTWARE

FUNDING · WORKSTATION CLOUD GPU CREDITS ITERATIVE COMPUTE

Project *site*
Paper · data · code
▶ weathon.github.io

Wenqi — *Fall 2026*
Industry roles · PhD positions

Qingyun — *on the market*
Looking for jobs · research Master's / PhD