



UNIVERSAL AESTHETIC ALIGNMENT NARROWS ARTISTIC EXPRESSION

Wenqi Marshall Guo, Qingyun Qian, Khalad Hasan, Shan Du


University of British Columbia (Okanagan) · Department of CMPS · Weathon Software, Canada

Image generators quietly **sanitize** the prompts they should follow, and reward models **punish** the answer the user actually asked for.

i Art the reward model rejects

HPSv3 TYP. 0–15


Two canonical paintings + three LAPIS works, all in **negative** reward territory. HPSv3 cannot distinguish deliberate aesthetic deviation from generation failure.




The Scream · Munch 1893
HPSv3 **5.23**




Luxe, Calme et Volupté · Matisse 1904
HPSv3 **1.73**



–4.28
landscape



–3.21
still life




–1.25
portrait

ii Emotion bias · **toxic** positivity

PROMPT-FAITHFUL = PUNISHED

Same face edited into three negative emotions. HPSv3 — asked to find the negative one — still picks happy. DanceFlux, asked to generate negative emotion, returns neutral or happy.



Happy 14.60 Sad 12.80 Anger 11.40 Fear 10.60

Reward-model accuracy on negative emotion

Model	Anger	Fear	Sad
BLIP (unaligned)	0.96	0.79	0.95
HPSv2	0.70	0.64	0.88
HPSv3	0.19	0.32	0.44
ImageReward	0.55	0.49	0.77

81%

HPSv3 picks happy when the prompt asks for anger.

iii DanceFlux vs Flux Krea on **normal** anti-aesthetic P_a









SUGAR-WATER DEFAULT

v Reversed alignment

TWO LOOPS

Each P_a explicitly asks for blur, deep shadow, melted shapes, disharmony, or chaotic composition. **Flux Krea** renders it.



DanceFlux returns hyper-saturated Pinterest-grade outputs anyway. LLM-judge anti-aesthetic coverage: Krea » DanceFlux (**0%** on every sample).



iv Image New Speak · **DanceFlux vs Flux Krea**, same prompt

100-PAIR DATASET · SAME DIRECTION EVERY TIME

Same socially-critical prompts go to **DanceFlux** (aesthetic-aligned) and **Flux Krea** (same Flux family, narrow-aligned but faithful). DanceFlux sanitizes; Krea keeps the critique.



User aligned **to the model**

Private loop: user asks for blur, decay, or grief; model returns its candy-gloss default and implicitly teaches that *this* is what good output looks like. The user adjusts.

Public loop: polished outputs flood feeds. Audiences internalize the narrow vocabulary as the default; new preference data folds this back into the next aligned model. The loop tightens — risking a cultural **mode collapse** in art.

And the suppressed content is not unsafe — it's critical, abstract, or emotionally negative imagery. Pre-emptive sanitization protects *corporate reputation*, not users.

“Rather, in the ugly, art must denounce the world that creates and reproduces the ugly in its own image.”
— T. W. Adorno, *Aesthetic Theory* (1984)

10 gen models · 7 reward models · 300 prompts · 2,928 real photos · $p < 10^{-10}$ for DanceFlux · HPSv3 prefers clean-but-wrong by +5.9 pts.

Why it matters · **value capture**

Reducing aesthetics to one reward score is *value capture* — the goal shifts from **make aesthetic images** to **make images that score high**. Pre-emptive exclusion of non-mainstream outputs is **pre-emptive governance**: it removes the user's capacity to dissent.

Call: alignment that exposes user-controllable preference strength, draws on diverse annotators, and stays transparent about what it optimizes.

Real **anti-aesthetic** photographs


From **aas_real_images** — intentional artistic choices that HPSv3 / ImageReward still rank deep in the **negative**. Score shown: ImageReward.



Project **site**
Paper · data · code
weathon.github.io



Wenqi — **Fall 2026**
Industry roles · PhD positions



Qingyun — **on the market**
Looking for jobs · research Master's / PhD

Made possible by

INNOVATION
Canada Foundation for Innovation
FUNDING · WORKSTATION

Lambda
CLOUD GPU CREDITS

Weathon
SOFTWARE
ITERATIVE COMPUTE